Benchmark progress: crypto-autoresearcher

From bench.record/v1 records in Git; raw data in S3 by SHA-256. Derived numbers are labelled derived.

Records: open 0, measured 3, invalid 3, censored 0. Headline unit S: group operations / sqrt(n); the ECDLP cost unit, every phase priced.

Series in unit count: gain vs first record (derived, %)

Series in unit fmul: gain vs first record (derived, %)

Series summary

problem • metric	unit	first	latest	gain vs first	value/ref	pts
ic:prime-field-fresh-holdouts • distinct_targets	count	8	8	?	?	1
ic:prime-field-fresh-holdouts • distinct_targets	count	15	15	?	?	1
ic:prime-field-fresh-holdouts • source_construct	fmul	98,566	98,566	?	?	1
ic:prime-field-fresh-holdouts • source_queries	count	49,152	49,152	?	?	1
ic:prime-field-source-budgets • distinct_targets	count	10	10	?	?	1
ic:prime-field-source-budgets • first_checked_in	count	96	96	?	?	1
ic:prime-field-source-budgets • first_checked_jo	count	48	48	?	?	1
ic:prime-field-source-budgets • source_queries	count	98,304	98,304	?	?	1
ic:prime-field-source-control • distinct_targets	count	8	8	?	?	1
ic:prime-field-source-control • source_queries	count	49,152	49,152	?	?	1

Every value is quoted from the record it names, in that record’s unit. Percentages and ratios under ’derived’ are computed here from those values and are not measurements. Series in different units are never compared.