Formal target provenance inventory at a named main commit

Static evidence review

25 October 2026

Named main: `0b38ac1157f777cb4a77d35e53d7ed583c541eb4`.
Formal tree: `96f334e856c87d397b5db720899d833a75db0352`.
Scope: Repository blob and declared-path inventory plus archived hosted Lean receipts. This is no experimental replay or semantic sign-off.

Observations

Target YAML specs	87
Existing referenced theorem paths	87
Distinct declared source paths	30
EXP-prefix directories at named main	2,781
Directory IDs with exact target/source token pointer	5
Formal targets with such a pointer	17
Specs with canonical EXP identifier	12
Source paths under experiments/	6
Direct /runs/ source paths	0
Specs with explicit run reference	0
Exact hosted machine-verified receipts	87

The canonical IDs in those twelve specs are `EXP-BINSTD-8bdf86` and `EXP-FROB-91ee9c`. Source status labels are COMPUTED (54), PROVED (23), STATED (8), DERIVATION (1), and MEASURED WITH ALGEBRAIC EXPLANATION (1). These labels are not independently checked outcomes. The hosted Lean run `40168163325` completed successfully at main `5998553341fb29d3b9ca1a4c9ac6e58d96ca3ffa`. Its 87 archived receipts report machine verification, successful builds and axiom audits, no forbidden constructs, and semantic review pending. The full formal tree and all 30 declared source blobs equal the named main snapshot.

The reverse pointer map includes legacy numeric experiment IDs and finds a fifth referenced directory, `EXP-NCP-001`, omitted by the earlier six-hex-digit scan. The other 2,776 directory IDs have no exact token pointer in these target/spec-source pairs. This does not show that a generic theorem is inapplicable to them or that the five pointed experiments have semantic sign-off.

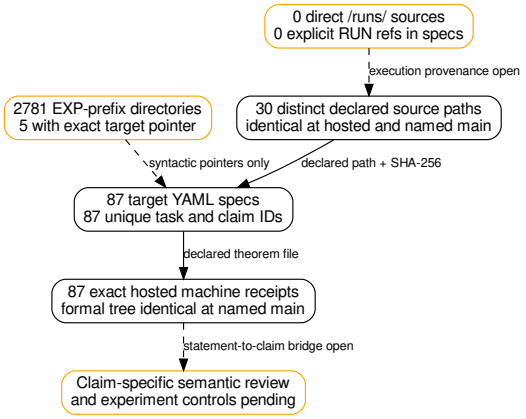

Solid links represent declared repository paths. Dashed links mark open execution provenance and statement-to-claim obligations. Counts describe the named commit only.

Interpretation and next check

The earlier local per-target sweep covered 84 targets. The later hosted 87-target run closes the machine-verification gap for the identical named formal tree. No independent semantic review across all claims is established. Machine acceptance cannot show that a theorem matches a historical experiment. Absence of a direct run path does not prove that no other provenance exists. Generic theorems may legitimately lack experiment IDs.

Before a claim is promoted, review its frozen source and run evidence, precise theorem hypotheses, model bridge, controls, and independent semantic assessment. Empirical measurements require provenance and uncertainty checks even when no theorem is suitable. No quantitative performance or attack-cost claim is made here. The open repository problem is `problems/formal-target-semantic-fidelity-unreviewed.md`. Existing scientific graphs and measured dashboards are unchanged because no new result or decision was recorded.