Benchmark dashboard provenance audit

Nine archived N13 rows, checked 2–3 November 2026

Checked publication. The public crypto-autoresearcher Pages build `6d5e7c5` served the pinned nine-row cryptanalysis snapshot exactly. The live comparison payload's benchmark object equals the committed JSON; all nine receipt selectors and the catalog name its SHA-256 `4cc3b31a...568ae5e`. Pages run 41512121316 succeeded, and the captured live meta payload names the same build commit. A second successful Pages run 41653395285 deployed current main `9b721d7` on 3 November; its captured comparison payload is byte-identical to the first readback.

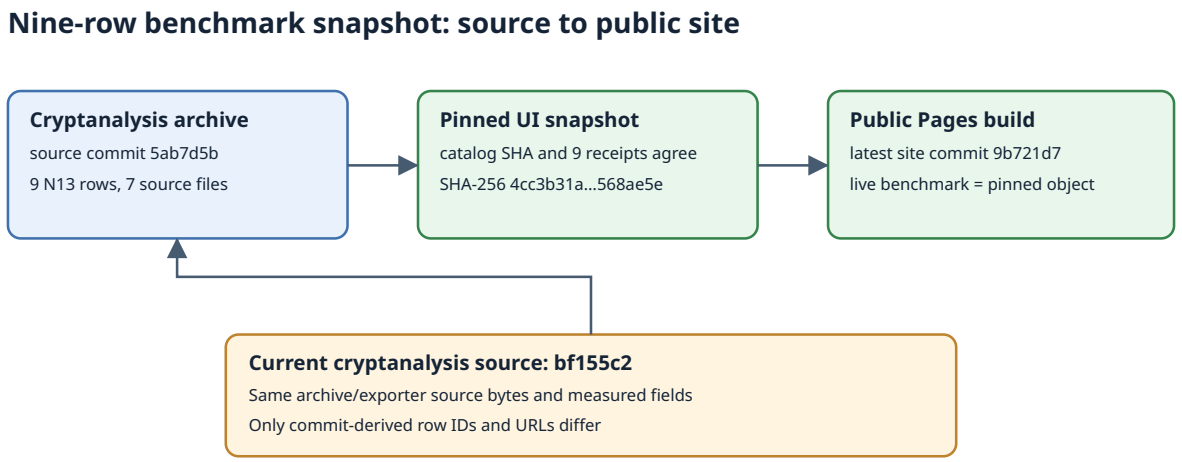

Current source check. The stored snapshot pins cryptanalysis `5ab7d5b`, the archived revision specified by the UI contract. A clean export at cryptanalysis `bf155c2` returned nine rows and seven source files. After removing only commit-derived source labels, row IDs and URLs, the two exports are equal: zero measurement, curve-identity, run-ID or source hash differences. The primary archive, candidate and workload manifests, exporter, and curve-identity tool have no Git diff between those commits. Updating the pin would alter labels without changing the data.

Reproduction and scope. The adjacent `check.py` compares the committed snapshot and receipt registry with archived current-export and live-site payloads; `check-result.json` records its passing output and input digests. This checks source and publication fidelity for the nine N13 toy rows. It does not rerun the archived experiment or certify its recorded runtime measurements. Other dashboard panels and experiment families are separate audit targets. No performance graph changed because its numeric inputs and their source hashes did not change.