Dashboard corpus integrity baseline

Published snapshot `e23868a6e5` · 2026-11-07 UTC · Conductor T-1870

The published dashboard and an exact-commit local build have identical `data/integrity.json` bytes (SHA-256 `805dbb6acbf7328baa2ca9df7c78748e949da75120fa35a874f6c0e8eb337f9d`). The local build indexed 29,068 records and completed its exact parse sweep.

Diagnostic	Count	Reading
Malformed YAML	24	Exact committed files fail parsing
Duplicate IDs	258	Stored copies and collisions mixed
ID/path mismatches	358	Diagnostic, not an invalidity verdict
Dangling IDs	1,048	Unresolved by the dashboard index

Published dashboard integrity: raw diagnostics and review boundary

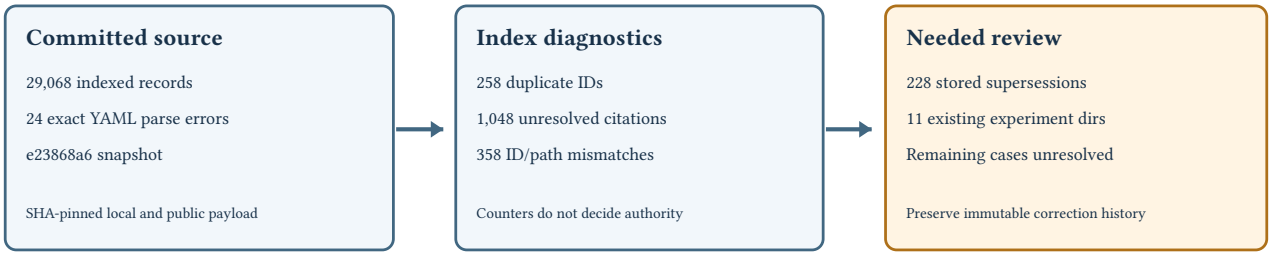

A dashboard diagnostic is not an experiment verdict or a reason to edit historical records in place.

The read-only audit script rebuilt the index, asserted all counts, and retained full duplicate-path hashes and all 1,048 unresolved IDs in `baseline.json`. The public payload displays only the top 500 dangling IDs. No ledger record was changed.

Classification and next checks

Of 258 duplicate groups, 228 contain stored schema supersessions, 12 contain archive copies, 12 contain other correction copies, five are active-ledger or legacy-root pairs, and one is an experiment header fallback collision. Only one group has byte-identical files; the other 257 require authority-aware reading. A stored correction is not interchangeable with its predecessor.

Eleven dangling `EXP-*` IDs have matching experiment directories. Two use nested IDs in their specifications; one specification fails YAML parsing; eight directories have no specification at that path. The remaining 1,037 IDs have no such direct experiment-directory evidence and need record-by-record custody review. Neither the dangling count nor a directory proves an experiment was executed or is missing.

All 24 malformed paths failed an independent exact YAML parse and have source hashes in `baseline.json`. They are current exact-detail coverage defects. The dashboard reports them visibly. Review immutable correction paths before any source change; inspect the 18 duplicate groups outside the explicit schema-supersession/archive classes and classify dangling IDs against planned work and correction inventories.

The nine primary IC benchmark rows and their graph values were independently audited in PR 3165 This corpus-wide integrity review adds no performance claim and changes no historical run record.

Reproduce with `audit.py` at commit `e23868a6e5`; see `report.md` for the full command and interpretation.