Run report: EXP-ICPERF-127dbd

This report records what ran and what it measured. It is generated by tools/runbook.py from the session journal. It states no verdict, support, weakening or significance; the evidence review does that.

- Runbook: experiments/EXP-ICPERF-127dbd/runbook.yaml (sha256 b3676f85f7700962619108a7a93ec8683213612fce3d211c00f85f09a245dd5b)
- Session: experiments/EXP-ICPERF-127dbd/runbook-sessions/RBS-ICPERF-de8978
- Journal tip: 2ed8c747ca282963754a834ddd8a5f01965672b22afeaba7d061067ac875d9e0 (8 entries)
- Commit: 4b285cb838df9b55b4a6efcdab6a66a6a9a40d02; tracked files dirty: False
- Started 2026-11-06T23:09:36.792423Z, finished 2026-11-06T23:09:46.120210Z; status stopped

Purpose

Executes protocol IC-C1-DEFERRED-GUARD-FP28-20261106 version 2 (protocol_sha256 9844f78d..., which carries AMENDMENT-1; the version 1 hash was 43b1ca91...). S1 and S2 are each submitted once through launch-remote by implementation/launch_eks.sh under the trial supervisor, which downloads the job prefix. check_job.py checks transport and identity of each download. The frozen analyze-c1.sh then evaluates the integrity, control and parity gates of both sessions. Observations only; no scientific outcome is asserted.

Steps

step	status	what it does
S01-probe	passed	Runs the bound probe_s01.sh once. In order it checks the first token of aws --version against aws-cli/1.46.2 (the CLI whose memory use was measured), prints aws sts get-caller-identity, checks the sha256 of the local pinned ecbench, recomputes the sha256 of request-s1.json and request-s2.json against protocol.json extensions.requests, checks the four launcher scripts against remote-pins.sha256, and prints the crypto checkout HEAD. It exits nonzero at the first mismatch, which stops the session before any launch. The step requires Account 104186811093 in its output; no AWS write call is made.
S02-trials	failed	Trial T1 runs launch_eks.sh on request-s1.json and trial T2 on request-s2.json under the trial supervisor, each followed by check_job.py. T2 starts only if T1 is output_validated. Controls is empty because the C1 controls (incumbent-aa, rho-neg-aa, pc-fb2x) are arms inside each ecbench job, not plan trials; their gates are evaluated by the frozen analyze-c1.sh in step S03-verify. No step retries a launch. The declared artifacts are the files under 1 MiB; records.jsonl and plan.json of each session are bound through job/artifacts/SHA256SUMS and job_files in raw-result.json.
S03-verify	not_run	Runs analyze-c1.sh once on session 1 and session 2 with the pinned ecbench; the validity assertions cover only the integrity, control and parity gate fields of verdict.json, never a claim outcome.

Units

0 valid of 2; invalid: ['T1', 'T2']

trial	run	receipt status	valid
T1	RUN-ICPERF-8b0d47	not_run	False
T2	RUN-ICPERF-5ffa79	not_run	False

Metrics

Each value is computed by tools/runbook.py from the recorded outputs named in its source. Labels say whether a value was measured or derived from measured values.

metric	label	unit	n	values
launch_remote_exit	measured	exit_status	0	count=0, max=null

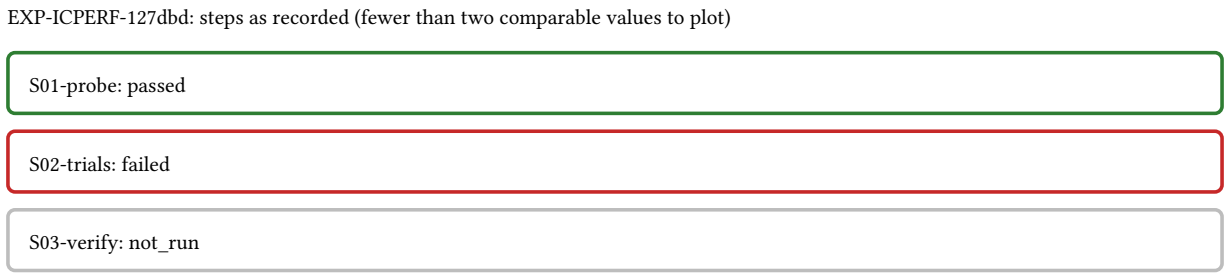

Figure 1: Recorded values from this session only; no new computation.

Validity, proofs and formal obligations

- Execution validity: incomplete
- Control validity: not_applicable
- Inferential validity: not_assessed: review-evidence and the Coordinator decide
- Execution proofs: {"chain": "incomplete: docs/execution-proof-chain.md milestones are open", "failed_units": 0, "unavailable_units": 0, "verified_units": 0}
- Audit: not audited: tools/runbook.py audit (audit-runs skill)

Where it stopped

S02-trials: infrastructure_error: claim failed (exit 2); see experiments/EXP-ICPERF-127dbd/runbook-sessions/RBS-ICPERF-de8978/steps/S02-trials.attempt-1.claim.stderr

A stop is an operational fact. It is not evidence about the hypothesis.