Run report: EXP-ICPERF-283a70

This report records what ran and what it measured. It is generated by tools/runbook.py from the session journal. It states no verdict, support, weakening or significance; the evidence review does that.

- Runbook: experiments/EXP-ICPERF-283a70/runbook.yaml (sha256 2d5aa752d5f2b2fce495b9a836d8637bb6ffed0b00ff3b97c5a8f63396149840)
- Session: experiments/EXP-ICPERF-283a70/runbook-sessions/RBS-ICPERF-67df7d
- Journal tip: c97b6c9d4a64ccbaaf9afbd35c733efba13d22ee6efc26cf56b961dcdb4df7cf (9 entries)
- Commit: af94114e9c2ff7fac8c7aeedf588c648c2098b99; tracked files dirty: False
- Started 2026-10-30T15:23:35.562921Z, finished 2026-10-30T15:25:13.248913Z; status completed

Purpose

Three deterministic trials. The primary trial draws 64 geometric counts at success probability 1/50 from seed 20261017 and records the sample moments. Two wrong samplers run beside it. The assertions decide the frozen success criterion by exact match with 2943, 289567 and 28; the band flags are descriptive. Observations only. Host latency is not an input.

Steps

step	status	what it does
S01-probe-python	passed	Writes the python3 version, implementation and executable path to environment.json in this step output directory.
S02-trials	passed	Runs run.py then check.py for each frozen trial in its own write-once directory; the assertions record the success criterion and each control.

Units

3 valid of 3; invalid: none

trial	run	receipt status	valid
geometric-probe-sample	RUN-ICPERF-f21eb1	output_validated	True
ctl-neg-wrong-triple (control)	RUN-ICPERF-0ef41e	output_validated	True
ctl-neg-support-zero (control)	RUN-ICPERF-74e9a2	output_validated	True

Metrics

Each value is computed by tools/runbook.py from the recorded outputs named in its source. Labels say whether a value was measured or derived from measured values.

metric	label	unit	n	values
cv_in_band	measured	fraction	1	fraction_true=1
elapsed_seconds	measured	seconds	1	max=0.000653293
exact_match	measured	fraction	1	fraction_true=1
median	measured	count	1	max=28, min=28
median_ratio_in_band	measured	fraction	1	fraction_true=1
sample_sum	measured	count	1	max=2943, min=2943
sum_squares	measured	count	1	max=289567, min=289567

EXP-ICPERF-283a70: steps as recorded (fewer than two comparable values to plot)

S01-probe-python: passed

S02-trials: passed

Figure 1: Recorded values from this session only; no new computation.

Validity, proofs and formal obligations

- Execution validity: complete
- Control validity: passed
- Inferential validity: not_assessed: review-evidence and the Coordinator decide
- Execution proofs: {"chain": "incomplete: docs/execution-proof-chain.md milestones are open", "failed_units": 0, "unavailable_units": 0, "verified_units": 0}
- Audit: not audited: tools/runbook.py audit (audit-runs skill)