Run report: EXP-ICPERF-284bc4

This report records what ran and what it measured. It is generated by tools/runbook.py from the session journal. It states no verdict, support, weakening or significance; the evidence review does that.

- Runbook: experiments/EXP-ICPERF-284bc4/runbook.yaml (sha256 e5850bc6a3076225edb948c931884fb1439383e74109ef8e7a36af3dfbe1a62c)
- Session: experiments/EXP-ICPERF-284bc4/runbook-sessions/RBS-ICPERF-8cdc92
- Journal tip: 534b1a29241d91679683028f5afe99a2bf52d1a3483826a983df48833a1ccdff (12 entries)
- Commit: 9d13fcfde5dfb40c5f34dc427a8c97e08cde7108; tracked files dirty: False
- Started 2026-10-30T15:16:40.897178Z, finished 2026-10-30T15:20:25.511083Z; status completed

Purpose

Six deterministic trials. The primary trial computes the reduced Bernstein-Lange price ratio and the online ratios at three frozen (r, A) cells. Three arithmetic-consistency checks and two planted negative controls run beside it. The assertions decide the frozen success criterion (925, 4, true, true). Observations only. The n=83 host arm is not an input.

Steps

step	status	what it does
S01-probe-python	passed	Writes the python3 version, implementation and executable path to environment.json in this step output directory.
S02-trials	passed	Runs run.py then check.py for each frozen trial in its own write-once directory; the assertions record the success criterion and each control.

Units

6 valid of 6; invalid: none

trial	run	receipt status	valid
price-ratio-identity	RUN-ICPERF-8e3acd	output_validated	True
ctl-consistency-a-change (control)	RUN-ICPERF-2d715d	output_validated	True
ctl-consistency-r-change (control)	RUN-ICPERF-2c627f	output_validated	True
ctl-bsgs-enumeration (control)	RUN-ICPERF-779d41	output_validated	True
ctl-neg-perturbed-price (control)	RUN-ICPERF-30596e	output_validated	True
ctl-neg-precomputation-model (control)	RUN-ICPERF-850008	output_validated	True

Metrics

Each value is computed by tools/runbook.py from the recorded outputs named in its source. Labels say whether a value was measured or derived from measured values.

metric	label	unit	n	values
all_match_bsgs	measured	fraction	1	fraction_true=1
elapsed_seconds	measured	seconds	1	max=0.0157668
matches_preregistered	measured	fraction	1	fraction_true=1
price_ratio_den	measured	count	1	max=4, min=4
price_ratio_num	measured	count	1	max=925, min=925
ratios_agree	measured	fraction	1	fraction_true=1

EXP-ICPERF-284bc4: steps as recorded (fewer than two comparable values to plot)

S01-probe-python: passed

S02-trials: passed

Figure 1: Recorded values from this session only; no new computation.

Validity, proofs and formal obligations

- Execution validity: complete
- Control validity: passed
- Inferential validity: not_assessed: review-evidence and the Coordinator decide
- Execution proofs: {"chain": "incomplete: docs/execution-proof-chain.md milestones are open", "failed_units": 0, "unavailable_units": 0, "verified_units": 0}
- Audit: not audited: tools/runbook.py audit (audit-runs skill)