Run report: EXP-SEMBIN-cb5e25

This report records what ran and what it measured. It is generated by tools/runbook.py from the session journal. It states no verdict, support, weakening or significance; the evidence review does that.

- Runbook: experiments/EXP-SEMBIN-cb5e25/runbook.yaml (sha256 054b45287bd5e34006dfc07f61af4b0a7929e59d3279bd37fc206cf36b60a7f0)
- Session: experiments/EXP-SEMBIN-cb5e25/runbook-sessions/RBS-SEMBIN-19ce78
- Journal tip: b445c899f15bfb704538e8fbbe0b3b7951538b4c1fcd50181c1607b71bc973bf (4 entries)
- Commit: 9626af41d014ccdc79445b2144f8e171df5a220d; tracked files dirty: False
- Started 2026-11-08T21:53:15.583327Z, finished 2026-11-08T21:53:15.623188Z; status stopped

Purpose

Counterexample search against the degree-4 mutant closure of coset-shifted EQS4 systems (Nagao 2015/984 binary chain). The run writes the closure records M-1 to M-10, the controls and the certificate records into one run directory and aggregates them with the pre-registered analysis script. Observations only: the run states no verdict on the hypothesis.

Steps

step	status	what it does
S01-probe-host	failed	Writes environment.json in this step output directory with the free disk space in bytes at the experiment directory, the MemAvailable line of /proc/meminfo in kilobytes, the value of the CRYPTO_AR_BLOBSTORE environment variable, and whether the Sage executable named by fixture F-6 exists.
S02-blobstore-round-trip	not_run	Runs store_from_env().put and .get of tools/artifact_store.py on a fixed 40 byte string and writes blob-roundtrip.json in this step output directory with the object address, the sha256, the length and whether the bytes read back are equal to the bytes stored.
S03-create-run-dir	not_run	Creates experiments/EXP-SEMBIN-cb5e25/runs/RUN-SEMBIN-d1204f and fails if it already exists.
S04-copy-collision-search	not_run	Copies experiments/EXP-SEMBIN-cb5e25/collision-search.json to experiments/EXP-SEMBIN-cb5e25/runs/RUN-SEMBIN-d1204f/collision-search.json without overwriting, and writes collision-copy.json in this step output directory with the sha256 of the source and of the copy.
S05-fixtures	not_run	Runs code/fixtures.py with the protocol hash and writes fixtures.json into the run directory. It exits 0 only when every fixture passes (SR-1). The time of this step is outside the 43200 second total of SR-8 and is recorded in the session journal.
S06-fixtures-record	not_run	Binds the fixture verdicts to the journal by digest and asserts that F-1 to F-8 each report pass.
S07-trials	not_run	Runs each code/adapter.py invocation of the frozen plan in its own supervisor run directory, followed by code/check_trial.py. All of them write their records into the one run directory of step S03. The 43200 second total of SR-8 is charged from the first trial.
S08-analysis	not_run	Writes controls.json and summary.json into the run directory and prints the validity flags.
S09-analysis-record	not_run	Binds the analysis outputs to the journal by digest and asserts that the SR-4 and SR-5 flags are false.

Units

0 valid of 0; invalid: none

trial	run	receipt status	valid

Metrics

Each value is computed by tools/runbook.py from the recorded outputs named in its source. Labels say whether a value was measured or derived from measured values.

metric	label	unit	n	values
closure_censored_fraction	measured	fraction	0	count=0, fraction_true=null
trial_wall_seconds	measured	seconds	0	count=0, max=null, median=null, sum=null

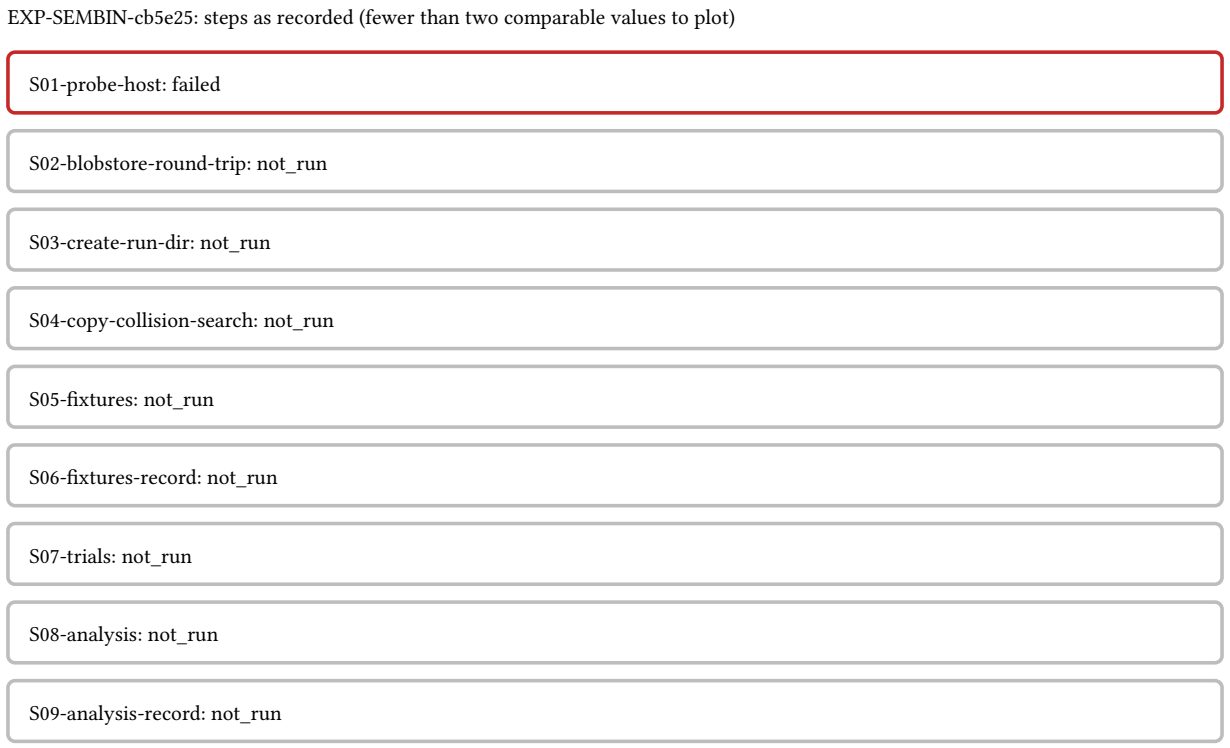

Figure 1: Recorded values from this session only; no new computation.

Validity, proofs and formal obligations

- Execution validity: incomplete
- Control validity: not_applicable
- Inferential validity: not_assessed: review-evidence and the Coordinator decide
- Execution proofs: {"chain": "incomplete: docs/execution-proof-chain.md milestones are open", "failed_units": 0, "unavailable_units": 0, "verified_units": 0}
- Audit: not audited: tools/runbook.py audit (audit-runs skill)

Where it stopped

S01-probe-host: invalid_measurement: /blobstore matches "^(s3://|file:///)": observed null; /sage equals true: observed false

A stop is an operational fact. It is not evidence about the hypothesis.